

Denser Perceptron Networks

Thesis Report

Student: Sarosh Krishan

Supervisor: Thanh Phuong Nguyen

Introduction

Perceptrons (hereafter referred to as "neurons") are fundamental units in neural networks capable of learning input features. Each neuron assigns learnable weights (parameters) to its input variables and generates a numerical output, which represents the learned feature for a specific input. Neurons can be classified into three types: **input neurons**, which connect directly to the input data; **intermediate neurons**, which process outputs from other neurons and learn more complex features based on these combinations; and **output neurons**, which provide the final output of the network.

In traditional **Multi-layer Perceptron (MLP)** networks, neurons are grouped into layers, allowing for parallel processing of multiple neurons. This layer-based structure enables the network to learn and process multiple features from a given input simultaneously. The advantage of layers is that they enhance the network's adaptability by increasing the number of learnable features while keeping the overall time complexity manageable. Without layers, the network would consist of neurons arranged sequentially, with each neuron learning a single feature, which may become inefficient as the complexity of the data increases.

In practice, MLPs aim to construct a network with numerous connections, modifying these connections during training to optimize the output. However, research suggests that many of these connections are redundant. In fact, studies indicate that up to 90% of all connections in a neural network are not significantly contributing to its performance. This finding implies that a smaller, carefully selected subnetwork could perform similarly to the original, fully connected network. Thus, MLPs attempt to explore all potential connections and retain those that offer the most value for the desired outcome.

Despite this, traditional MLPs intentionally omit certain types of connections to improve efficiency. Notably, neurons within a single layer do not interact with each other, limiting the network's capacity to learn complex, interdependent features. Additionally, neurons separated by more than one layer lack direct connections, relying on intermediate neurons to propagate information between them. Although skip connections—where the input to a layer is added to its output—partially address this issue, they come with three key limitations:

1. Skip connections generally only span a single layer, and neurons more than two layers apart remain disconnected.
2. Skip connections require the input and output of a layer to have the same dimensionality, which restricts the flexibility of network design.
3. While skip connections combine inputs, they do not allow for individual weights to be assigned to each input, thus reducing the network's ability to learn from them independently.

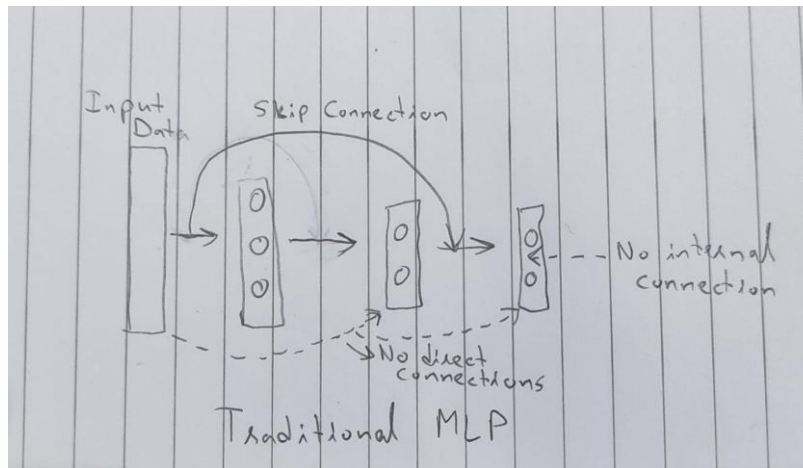


Figure 1 Limitations of Traditional MLPs

Although space considerations played a significant role in the original design of neural networks, modern hardware is more than capable of handling much larger networks. Despite this, traditional MLPs have proven sufficient for many tasks, and increasing the number of neurons per layer has often been used to leverage the added capacity of modern hardware. This raises the question: Can we optimize the use of space in neural network architectures to achieve greater efficiency?

Denser Perceptron Networks (DPNs)

This thesis introduces **Denser Perceptron Networks (DPNs)**, a framework designed to eliminate restrictions on neuron connectivity. The goal is to allow any two neurons to connect, forming a directed acyclic graph (DAG) like network. This framework seeks to enhance neural networks by improving their connectivity while minimizing the associated increases in time and space complexities.




Figure 2 An example of a DPN

In theory, DPNs treat neurons as individual objects that can process inputs (also referred to as a forward pass) independently. These neurons may still depend on either the input data or other neurons for input, but the framework enables greater flexibility in forming connections across the network.

Maximal DPNs

To assess the time complexity of DPNs, we first examine the densest possible network configuration, known as **Maximal DPNs**. In this arrangement, neurons are connected in a sequential manner, where the first neuron is connected only to the input, the second neuron is connected to both the input and the first neuron, and so on. This structure leads to a fully connected network, where each neuron is linked to all preceding neurons, resulting in a Maximal DPN.




Figure 3 Maximal DPN

1. Object-based Representation with Partial Inputs and Propagation.

One potential method to achieve the DPN framework is for neurons, as independent objects, to store partial inputs from sources that have completed their forward pass, while still waiting on outputs from other sources that haven't finished. Once a neuron receives all its necessary inputs, it performs its own forward pass and propagates its output to the subsequent neurons. However, this approach requires that the network be free of loops, as the presence of loops would result in deadlocks.




Figure 4 Functionality of a single Neuron in Object-based Representation

While this approach is conceptually simple and flexible, it is computationally expensive. Each neuron performs a storage and propagation step for every connection it has, leading to a large amount of redundancy. Even though the forward pass is only a single step once the inputs are complete, a worst-case scenario involves performing n forward passes, where n is the total number of neurons. Additionally, since each neuron stores its partial inputs, the outputs from a single neuron are duplicated for every neuron it propagates to. With each neuron also storing weights for every input, this method doubles the memory used compared to a traditional MLP network.

Time Complexity: $O(n^2)$, where n is the total number of neurons.

Space Complexity: $O(2d * n^2)$, where d is the feature size of the input data.

2. Sequential Representation with Compounding Input.

In this approach, we describe the densest form of a DPN network using a weight matrix, where each row corresponds to a neuron, and the columns represent the input features and the outputs from all preceding neurons, except for the final one. This results in a lower triangular matrix in a staircase shape.




Figure 5 DPN Matrix

Such a network contains all possible connections for a given number of neurons. Therefore, any other network with the same number of neurons is a subnetwork of this complete network. To process this network, neurons are processed sequentially, starting from the first row. The output of a neuron is appended to its input, providing the input for the subsequent neuron. By doing so, we eliminate the need to add partial inputs for each neuron individually. Instead, the same input can be compounded across the forward pass.

This approach assumes a maximum connection for every neuron. Disconnecting two neurons is achieved by zeroing out the weight value for the output neuron

corresponding to the input. This method resembles traditional MLP pruning, where weight values are zeroed out instead of being removed entirely.

While this method eliminates the propagation step from every neuron to its subsequent neurons, it still requires each neuron to temporarily store its input during the backpropagation process. Therefore, this approach does not reduce space complexity.

Time Complexity: $O(n)$.

Space Complexity: $O(2d * n^2)$.

3. Sequential Representation with Shared Input.

This approach is identical to the previous one, with the key difference being that instead of storing inputs temporarily, neurons index slices from an input variable shared across the entire network. This indexing does not add significantly to the time complexity but optimizes it by eliminating redundancy.

Time Complexity: $O(n)$.

Space Complexity: $O(d * n^2)$.

4. Block-based Representation with Partial Outputs and Shared Input.

When considering the lower triangular weight matrix and visualizing vertical lines at the edge of each neuron's weight vector, we see the formation of blocks within the matrix. The largest block contains the weights corresponding to the input features, and subsequent blocks contain the weights for the output of each neuron.




Figure 6 Blocks forming in a DPN Matrix

Rather than processing each neuron's weight vector individually, this approach processes the network one block at a time. The top-most neuron's forward pass is completed first, followed by calculating partial outputs for the remaining neurons. This method improves the time complexity by prioritizing the heaviest calculations (typically when the hardware is under lower load), and adds slight parallelization, as an input value is accessed only once, rather than repeatedly in a sequential approach.

This method is both energy and time-efficient, offering significant improvements over the previous approaches.

Time Complexity: $O(n)$.

Space Complexity: $O(d * n^2)$.

Cross Version Comparison

The following table demonstrates the performance of all four approaches on the same dataset (MNIST) with an identical neuron count (34).

Approach	Time Complexity	Total Training Time	Trainable Parameters	Test Accuracy
Object with Partial Inputs	$O(n^2)$	684.0s	466,480	0.9568
Sequential with Compound Input	$O(n)$	182.3s	466,480	0.9599
Sequential with Shared Input	$O(n)$	174.1s	466,480	0.9599
Block with Shared Input	$O(n)$	162.3s	466,480	0.9639

So far, we've created a theoretically optimized approach for maximal DPNs. However, maximal DPNs themselves are suboptimal neural network architectures due to the massive increase they cause to the time complexity compared to traditional MLPs. The following approaches aim to compromise on some of the enhancements of maximal DPNs for faster times, trying to find a best of both worlds approach.

Free Weights

When projecting traditional MLPs onto a lower triangular matrix, we notice weight blocks isolated both vertically and horizontally. These isolated weights represent the independent layers in traditional MLPs. The right side of the matrix defines the time complexity of the forward pass; the more layers present, the longer the processing time. However, the left side of the matrix contains **Free Weights**, zeroed-out weights that do not contribute to time complexity.

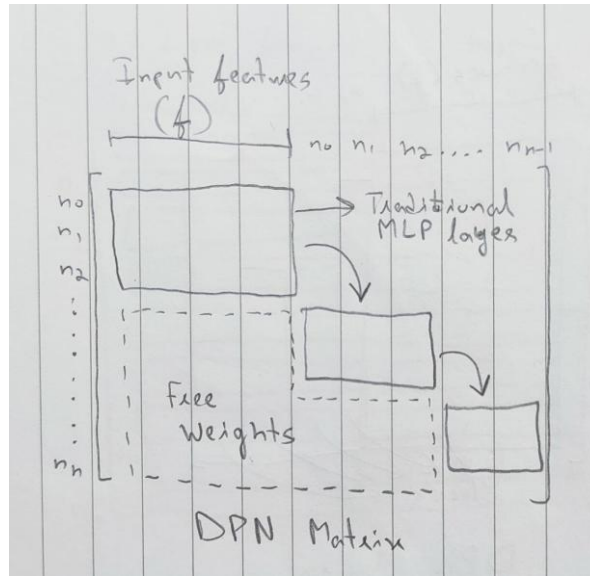


Figure 7 Free Weights in Traditional MLP

Using free weights in traditional MLPs is similar to concatenating the output of a layer to its input before passing it to the next layer. This method increases the space complexity of the MLP to $O(n^2)$ while keeping the time complexity at $O(l)$, where l is the number of layers. A major advantage of this approach is that subsequent layers can learn features based on both the input and the features from other previous layers, enabling the network to learn more complex patterns.

Minimal DPNs

While maximal DPNs maximize the number of connections in a perceptron network, they come with the trade-off of significantly increasing the time complexity from $O(l)$ (where l is the number of layers) to $O(n)$ (where n is the total number of neurons). Since $l \ll n$ in most perceptron networks, this added time complexity becomes a significant challenge.

On the other end of the spectrum, **Minimal DPNs** seek to minimize the number of blocks in the DPN framework. There are two possibilities:

1. If the total number of neurons is equal to the output size ($n == o$), only one block is needed, equivalent to a single MLP layer of size n .
2. If the total number of neurons exceeds the output size ($n > o$), the network must have at least two layers: one layer of size $n - o$ for the input, and a second layer of size o that takes both the input and the output of the first layer.

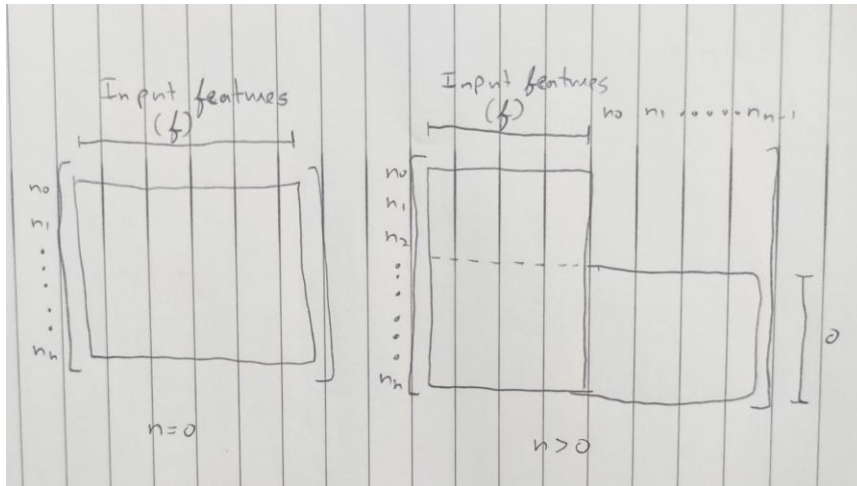


Figure 8 Minimal DPNs

Minimal DPNs provide the best time complexity of $O(1)$ and the best space complexity of $O(n^2)$.

Maximal DPNs with Pruning

The final approach to improving DPN efficiency is pruning, inspired by the **lottery ticket hypothesis**. According to this hypothesis, there exists a smaller subnetwork within a neural network that can perform similarly to the original network if trained on the same dataset. The process involves training the parent network briefly (e.g., for one epoch), pruning the network slightly, and then resetting the remaining weights to their original values before repeating the first step for a desired number of iterations, until the network reaches a desired size. Then, the best performing pruned version of the network is trained for the complete duration.




Figure 9 Maximal DPNs with Pruning

This approach is applied to Maximal DPNs by zeroing out weights along the right edge of the lower triangular matrix. This pruning allows some neurons to share the same right-side edges, enabling parallel processing. The time complexity of the network is improved to $\mathbf{O(l)}$, where \mathbf{l} is the number of vertical edges on the right side of the DPN matrix. Additionally, this method can reduce the space complexity by up to 90%, making it highly efficient.